

Amortized Analysis

Christian Wulff-Nilsen

Sixth lecture

Algorithms and Data Structures

DIKU

February 22, 2023

Overview for today

- Introduction
- Aggregate analysis
- The accounting method
- The potential method
- Applications:
 - Stacks with the MULTIPOP operation
 - Binary counters
 - Dynamic tables

Introduction

- So far in the course, we have looked at the worst-case running time/space of an algorithm
- Now, we consider instead the cost of maintaining a data structure under a sequence of operations
- For a stack, this would be a sequence of PUSH and POP operations
- Instead of bounding the worst-case time of a single operation, we wish to bound the *total worst-case time of all operations*
- This also allows us to *bound the average time of an operation*: simply divide the bound on the total cost by the number of operations
- We will study three methods to obtain such bounds on various types of data structures:
 - Aggregate analysis
 - The accounting method
 - The potential method

A stack with the MULTIPOP operation

- Consider a stack data structure with four operations:
 - $\text{PUSH}(S, x)$: pushes element x onto stack S
 - $\text{POP}(S)$: pops the top element from stack S (we assume S is non-empty just prior to the pop)
 - $\text{MULTIPOP}(S, k)$: pops the top $\min\{s, k\}$ elements from stack S where s is the number of elements on the stack just prior to the operation (we assume $\min\{s, k\} > 0$)
 - $\text{STACK-EMPTY}(S)$: outputs true if S is empty and false otherwise
- Consider an implementation where each call to PUSH, POP, and STACK-EMPTY takes $\Theta(1)$ worst-case time
- We implement MULTIPOP with $\min\{s, k\}$ successive POP operations, with calls to STACK-EMPTY to check if the stack becomes empty
- MULTIPOP runs in $\Theta(\min\{s, k\})$ worst-case time
- Technical detail: we require $\min\{s, k\} > 0$ as otherwise, $\Theta(\min\{s, k\}) = \Theta(0)$ is not a valid time bound; MULTIPOP spends $\Theta(1)$ worst-case time even when popping no elements

Worst-case running time of our stack data structure

- Consider n operations on an initially empty stack
- What is the worst-case *total* time of these operations?
- Worst-case time for a single operation is $O(n)$ (a single MULTIPOP operation can take up to $\Theta(n)$ worst-case time)
- Since there are n operations, we get a total worst-case time over all n operations of $O(n^2)$
- The average time per operation is thus $O(n)$
- Our analysis is correct but the $O(n)$ average bound is very weak
- Using amortized analysis, we can improve it to $O(1)$

Aggregate analysis applied to the stack example

- In aggregate analysis, we calculate an upper bound $T(n)$ on the total worst-case time of n operations and then calculate an upper bound on the average cost, or amortized cost, as $T(n)/n$
- Aggregate analysis is typically more refined than on the previous slide where we simply used the upper bound $O(n)$ for every operation
- Aggregate analysis for the stack example:
 - It suffices to bound the total worst-case time spent on PUSH and POP operations (Why?)
 - There are at most n PUSH operations in total
 - The number of POP operations (including those applied as part of MULTIPOP) cannot be larger than the number of PUSH operations
 - Hence, total worst-case time for all n operations is $O(n)$
 - The amortized cost per operation is $O(n)/n = O(1)$

Binary counter

- We now consider implementing a binary k -bit counter which starts at 0 and counts upwards with the operation INCREMENT
- The counter resets to 0 after 2^k increments (overflow)
- The counter is stored in an array $A[0 \dots k - 1]$ where $A[0]$ is the least and $A[k - 1]$ is the most significant bit

INCREMENT example

- Example of repeated application of INCREMENT with $k = 4$:

$A[3]$	$A[2]$	$A[1]$	$A[0]$
0	0	0	0
0	0	0	1
0	0	1	0
0	0	1	1
0	1	0	0
0	1	0	1
0	1	1	0
0	1	1	1
1	0	0	0
\vdots	\vdots	\vdots	\vdots

Implementation of INCREMENT

- INCREMENT is implemented as follows:

INCREMENT(A)

1 $i = 0$

2 while $i < A.length$ and $A[i] == 1$

3 $A[i] = 0$

4 $i = i + 1$

5 if $i < A.length$

6 $A[i] = 1$

$A[3]$	$A[2]$	$A[1]$	$A[0]$
0	1	1	1
1	0	0	0

A simple running time analysis

- Consider n INCREMENT(A) operations
- The worst-case time for a call to INCREMENT(A) is proportional to the number of bits that this operation flips which is $O(k)$
- Average time bound: $O(k)$
- We now strengthen this bound with aggregate analysis

Aggregate analysis applied to the binary counter example

- $A[0]$ flips for every call to $\text{INCREMENT}(A)$
- However, $A[1]$ only flips for every second call
- In general, $A[i]$ only flips for every 2^i th call, for $i = 0, \dots, k - 1$
- It follows that $A[i]$ only flips $\lfloor n/2^i \rfloor$ times in total
- The total number of flips over all n operations is thus

$$\sum_{i=0}^{k-1} \left\lfloor \frac{n}{2^i} \right\rfloor \leq \sum_{i=0}^{k-1} \frac{n}{2^i} = n \sum_{i=0}^{k-1} \frac{1}{2^i} < n \sum_{i=0}^{\infty} \frac{1}{2^i} = 2n$$



- The total worst-case time for all n operations is thus $O(n)$
- The amortized cost per operation is $O(1)$

The accounting method

- Consider a sequence of n operations to some data structure (for instance the stack or the binary counter)
- Let the time cost of the i th operation be $c_i, i = 1, \dots, n$
- In the accounting method, we assign **artificial costs** $\hat{c}_1, \dots, \hat{c}_n$ to the n operations
- \hat{c}_i is called the *amortized cost* of the i th operation
- Unlike aggregate analysis, we now allow different amortized costs to each of the n operations

The accounting method: requirement on amortized costs

- If $\hat{c}_i > c_i$, we overcharge the operation by the amount $\hat{c}_i - c_i$ which is stored as credit in specific objects of the data structure for later
- If $\hat{c}_i < c_i$, we undercharge the operation by the amount $c_i - \hat{c}_i$ and we pay for this difference with credit stored from previous operations
- Important requirement: for *any* sequence of n operations,


$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i$$


- Goal: obtain upper bound on $\sum_{i=1}^n \hat{c}_i$; by the above inequality, this will give an upper bound on $\sum_{i=1}^n c_i$
- The inequality says that, after having paid the actual cost $\sum_{i=1}^n c_i$ for the n operations, the remaining $\sum_{i=1}^n \hat{c}_i - \sum_{i=1}^n c_i$ credit stored in the data structure must be non-negative

The accounting method applied to the stack example

- Actual costs of stack operations (after scaling by a constant):
 - PUSH: 1
 - POP: 1
 - MULTIPOP: $\min\{s, k\}$
- We choose the following amortized costs:
 - PUSH: 2
 - POP: 0
 - MULTIPOP: 0
- The PUSH operation is overcharged by 1
- 1 of the 2 credits pays the actual cost of the PUSH operation and the remaining credit is left on the element that was pushed (imagine a coin left on the stack element)
- The POP operation is undercharged by 1 and its actual cost is paid for by the credit associated with the popped element
- The amount of credit associated with the stack is never negative


The accounting method applied to the stack example

- Example:

1 credit


The accounting method applied to the stack example

- Example:




The accounting method applied to the stack example

- Example:



The accounting method applied to the stack example

- Example:



The accounting method applied to the stack example

- For any sequence of n operations, this gives

$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i \leq 2n$$

- The **average cost per operation** is thus $O(1)$, which is what we obtained earlier using the aggregate method
- Similarly, we can use the accounting method for the binary counter example
- Note that the data structure does not keep track of credits; we only use them in the analysis

The potential method

- Similar to the accounting method except that credit is not stored with specific objects of the data structure (such as elements of the stack)
- Instead, credit is stored in a single place – “the bank”
- The current amount of credits in the bank is expressed by a *potential function* Φ
- Consider n operations to a data structure where D_0 is the data structure before the first operation and D_i is the data structure just after the i th operation, for $i = 1, \dots, n$
- For $i = 0, \dots, n$, we denote by $\Phi(D_i)$ the credit stored with the current data structure D_i

Amortized costs with the potential method

- For $i = 1, \dots, n$, if c_i is the actual cost of the i th operation, we define the amortized cost \hat{c}_i as

$$\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$$



- If $\Phi(D_i) - \Phi(D_{i-1}) > 0$, we overcharge the i th operation and put $\Phi(D_i) - \Phi(D_{i-1})$ credit into the bank
- If $\Phi(D_i) - \Phi(D_{i-1}) < 0$, we undercharge the i th operation and withdraw $\Phi(D_{i-1}) - \Phi(D_i)$ credit from the bank to help pay for the i th operation

Summing up amortized costs

- Recall from the accounting method the requirement that for any sequence of n operations,

$$\sum_{i=1}^n \hat{c}_i \geq \sum_{i=1}^n c_i$$

- We require the same for the potential method
- By a telescoping sums argument,

$$\sum_{i=1}^n \hat{c}_i = \sum_{i=1}^n (c_i + \Phi(D_i) - \Phi(D_{i-1})) = \left(\sum_{i=1}^n c_i \right) + \Phi(D_n) - \Phi(D_0)$$

Amortized cost as upper bound on actual cost

- By requiring that $\Phi(D_n) \geq \Phi(D_0)$, we get the desired inequality:

$$\sum_{i=1}^n \hat{c}_i = \left(\sum_{i=1}^n c_i \right) + \Phi(D_n) - \Phi(D_0) \geq \sum_{i=1}^n c_i$$


- Since we might not know n , we require $\Phi(D_i) \geq \Phi(D_0)$ for all $i \geq 0$
- In this case, we say that Φ is *valid*
- Typically, we pick Φ such that $\Phi(D_0) = 0$ and $\Phi(D_i) \geq 0$ for all $i \geq 0$
- This clearly ensures that Φ is valid
- In words, the amount of credit in the bank can never be negative

The potential method applied to the stack example

- Let D_0 be the initial stack and let D_i be the stack just after the i th operation
- We choose $\Phi(D_i)$ to be the number of elements on the stack D_i , for $i = 0, \dots, n$
- Φ satisfies the requirements on the previous slide since $\Phi(D_0) = 0$ and $\Phi(D_i) \geq 0$ for all i
- We thus have $\sum_{i=1}^n \hat{c}_i \geq \sum_{i=1}^n c_i$
- We now upper bound $\sum_{i=1}^n \hat{c}_i$ to get an upper bound on $\sum_{i=1}^n c_i$

Upper bounding $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$

- Let $i \in \{1, \dots, n\}$ be given and consider the i th operation
- Recall that $\Phi(D_i)$ is the number of elements on stack D_i
- If the i th operation is PUSH:

- $\Phi(D_i) - \Phi(D_{i-1}) = 1$ 
- Hence, $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 1 + 1 = 2$

- If it is POP:

- $\Phi(D_i) - \Phi(D_{i-1}) = -1$
- Hence, $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 1 - 1 = 0$

- If it is MULTIPOP:

- Let $k' > 0$ be the number of elements popped
- $\Phi(D_i) - \Phi(D_{i-1}) = -k'$
- Hence, $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = k' - k' = 0$

- In all cases: $\hat{c}_i \leq 2$

Upper bounding $\sum_{i=1}^n \hat{c}_i$ and hence $\sum_{i=1}^n c_i$

- We have shown that for each $i \in \{1, \dots, n\}$, $\hat{c}_i \leq 2$
- Hence, $\sum_{i=1}^n \hat{c}_i \leq 2n$
- This also upper bounds the total actual cost of the n operations:

$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i \leq 2n$$

- Hence, the average time spent per operation is $O(n)/n = O(1)$

Dynamic tables


- Consider some abstract data structure T which we refer to as a *table*
- It supports the insertion of an element with the operation INSERT and the deletion of an element with the operation DELETE
- We do not focus on the details of how T supports these operations; we only require that:
 - T is initially empty,
 - the memory used by T is allocated as an array of slots,
 - each operation is supported in $O(1)$ time, and
 - allocating/freeing space for an array of size k takes $O(k)$ time
- For simplicity of illustration, we assume in the following that T is a stack with INSERT corresponding to PUSH and DELETE corresponding to POP
- All of our observations immediately generalize to other data structures such as heaps and hash tables

Dynamic tables: Insertions only

- Assume for now that T only supports INSERT operations
- Goal: T should be able to dynamically allocate a new larger array once the old array is too small to contain all elements of T
- Using the potential method, we show how to do this using only $O(1)$ average time per insertion

Notation in the following

- num_i : number of elements in T just after the i th operation
- size_i : size of array associated with T just after the i th operation
- $\alpha_i = \text{num}_i / \text{size}_i$: the *load factor* α_i of T just after the i th operation
- If $\text{size}_i = 0$, define $\alpha_i = 1$
- α_i indicates how big a fraction of the array is filled with elements
- Example (non-empty entries of T are green):



$$\alpha_i = 3/4$$

Table expansion

- Initially, T has an empty array
- Just prior to inserting the i th element, if $\text{num}_{i-1} = \text{size}_{i-1}$ (equivalently, if $\alpha_{i-1} = 1$) then T is expanded:
 - A new array twice as big is allocated, i.e., $\text{size}_i = 2\text{size}_{i-1}$
 - The elements from the old array are copied to the new array
 - The old array is deallocated





Table expansion

- Initially, T has an empty array
- Just prior to inserting the i th element, if $\text{num}_{i-1} = \text{size}_{i-1}$ (equivalently, if $\alpha_{i-1} = 1$) then T is expanded:
 - A new array twice as big is allocated, i.e., $\text{size}_i = 2\text{size}_{i-1}$
 - The elements from the old array are copied to the new array
 - The old array is deallocated



The potential method applied to the dynamic table example


+

- Table expansion plus the insertion of the new element takes worst-case time $O(\text{num}_i)$ $5 + 1 = 6$
- Hence, if the i th operation requires a table expansion, we can set $c_i = \text{num}_i$ (after scaling by a constant factor)
- If no table expansion is required, $c_i = 1$
- We will now show how to bound the total cost $\sum_{i=1}^n c_i$ of a sequence of n operations using the potential method




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :



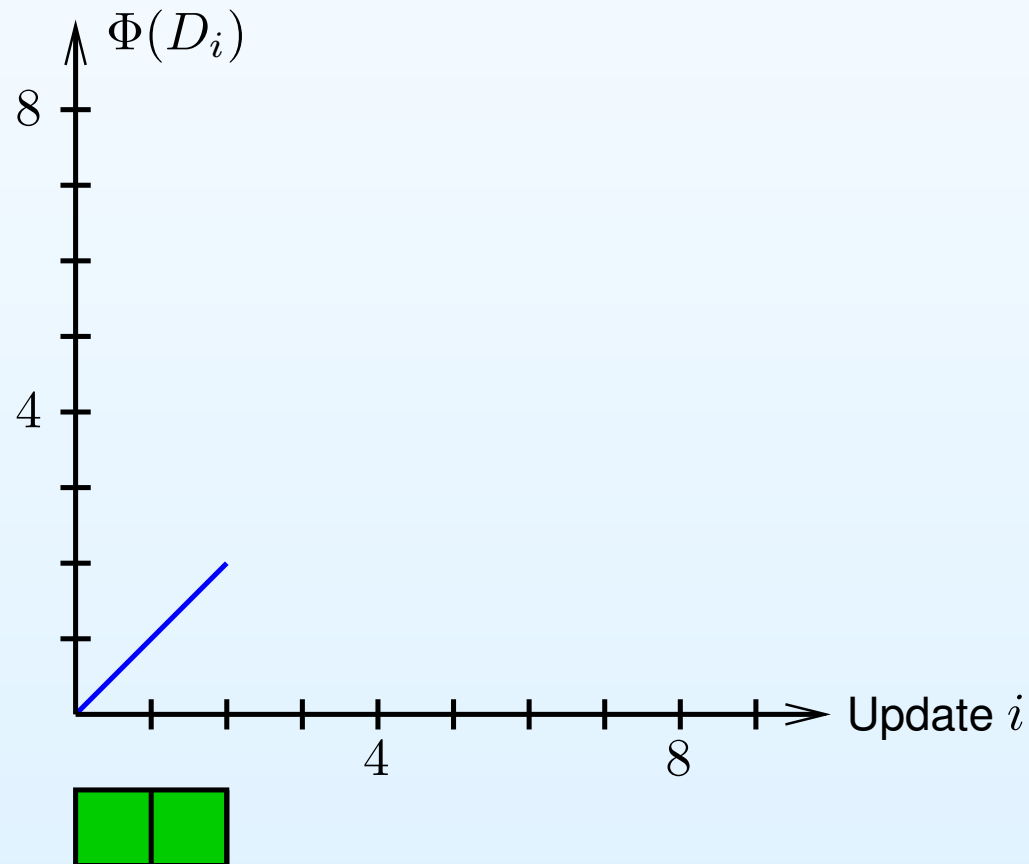
The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :



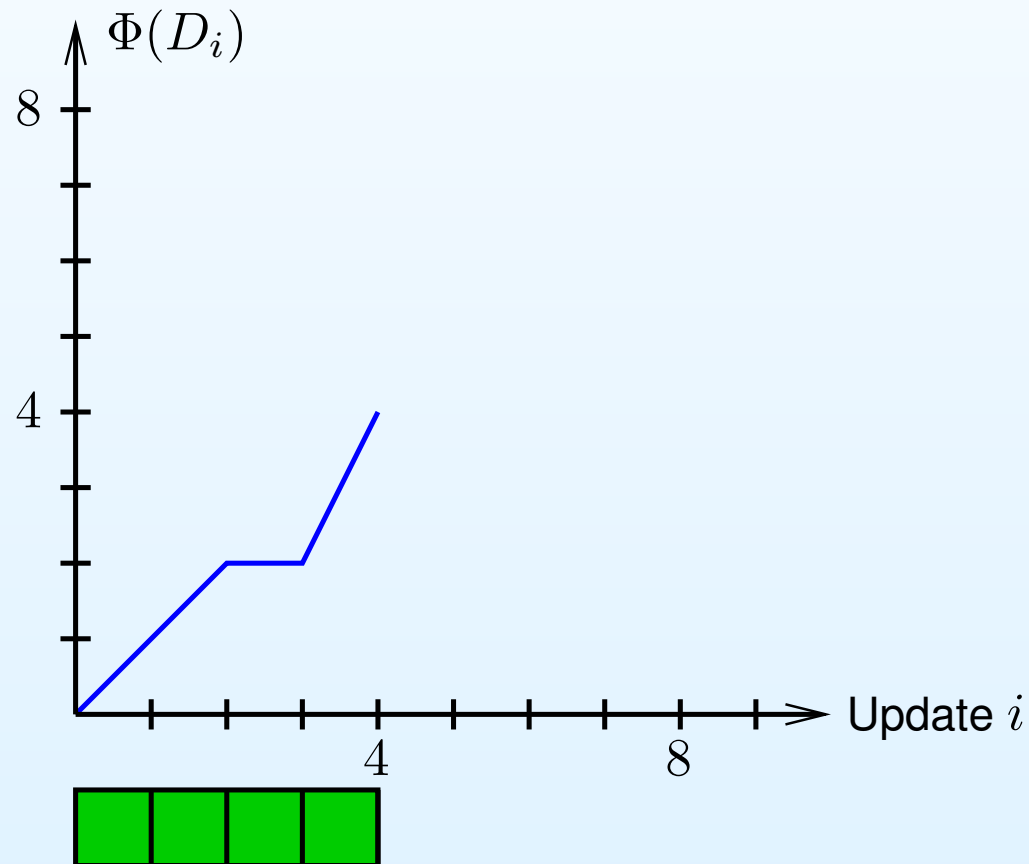
The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :



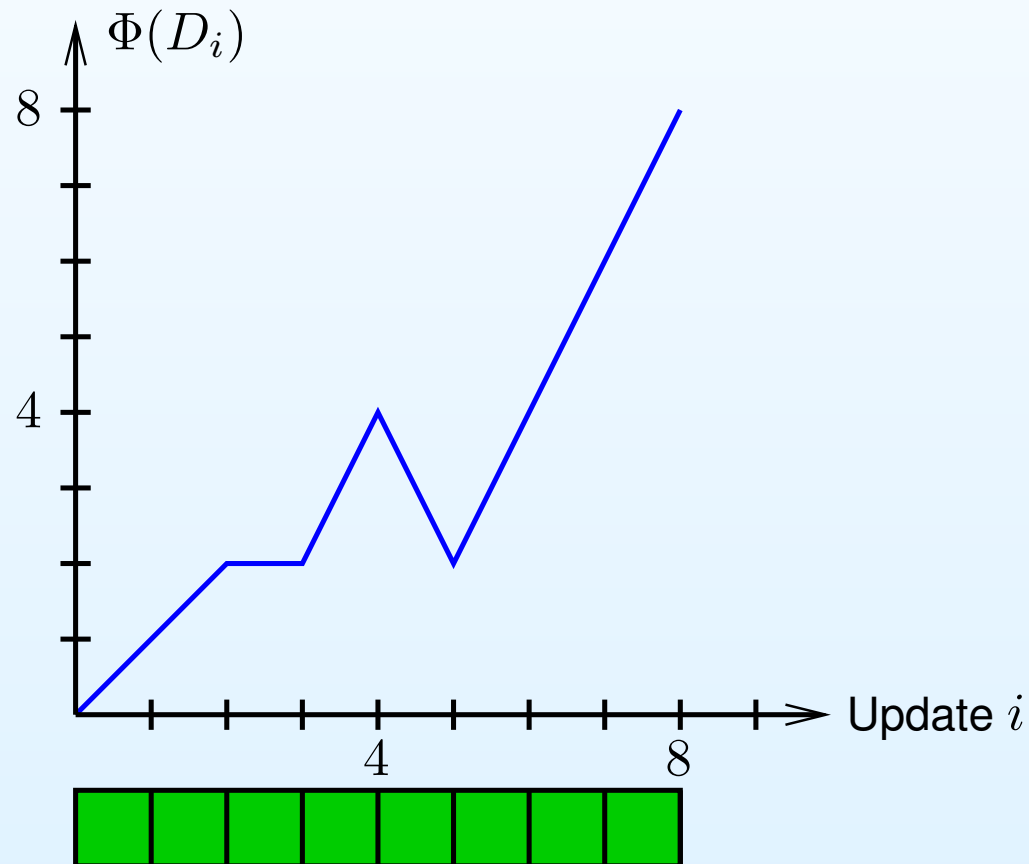
The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :




The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :



The potential method applied to the dynamic table example

- Let D_0 be the initial empty table T and for $i = 1, \dots, n$, let D_i be T just after the i th update
- We choose the potential function $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- Potential function as a function of i :



The potential function is valid

- We show that Φ is valid by proving that $\Phi(D_0) = 0$ and $\Phi(D_i) \geq 0$ for all i
- $\Phi(D_0) = 0$ is clear since T is initially empty
- Since T is always at least half full, $\Phi(D_i) \geq 0$ for all i
- It follows that Φ is valid and hence $\sum_{i=1}^n \hat{c}_i \geq \sum_{i=1}^n c_i$
- We can thus upper bound the total actual cost $\sum_{i=1}^n c_i$ by upper bounding $\sum_{i=1}^n \hat{c}_i$

The potential method applied to the dynamic table example

- Recall that $\Phi(D_i) = 2\text{num}_i - \text{size}_i$ for $i = 0, \dots, n$
- Consider the i th insertion operation, $i \in \{1, \dots, n\}$
- If no table expansion occurs, the amortized cost \hat{c}_i is

$$\begin{aligned}\hat{c}_i &= c_i + \Phi(D_i) - \Phi(D_{i-1}) \\ &= 1 + (2\text{num}_i - \text{size}_i) - (2\text{num}_{i-1} - \text{size}_{i-1}) \\ &= 1 + (2\cancel{\text{num}_i} - \cancel{\text{size}_i}) - (2(\cancel{\text{num}_i} - 1) - \cancel{\text{size}_i}) \\ &= 3\end{aligned}$$

1 + 2


The potential method applied to the dynamic table example

- Recall that $\Phi(D_i) = 2\text{num}_i - \text{size}_i$ for $i = 0, \dots, n$
- Consider the i th insertion operation, $i \in \{1, \dots, n\}$
- If table expansion does occur and $i = 1$, the amortized cost \hat{c}_i is

$$\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 2$$

The potential method applied to the dynamic table example

- Recall that $\Phi(D_i) = 2\text{num}_i - \text{size}_i$ for $i = 0, \dots, n$
- Consider the i th insertion operation, $i \in \{1, \dots, n\}$
- Now assume that table expansion occurs and $i > 1$:



- The amortized cost \hat{c}_i is

$$\begin{aligned}\hat{c}_i &= c_i + \Phi(D_i) - \Phi(D_{i-1}) \\ &= \text{num}_i + (2\text{num}_i - \text{size}_i) - (2\text{num}_{i-1} - \text{size}_{i-1}) \\ &= \text{num}_i + (2\text{num}_i - 2(\text{num}_i - 1)) - (2(\text{num}_i - 1) - (\text{num}_i - 1)) \\ &= 3\end{aligned}$$



Table contraction

- So far, we only considered the operation INSERT
- Now, we also include the operation DELETE
- To save memory, we want the array to contract to a smaller array whenever the load factor $\alpha_i = \text{num}_i / \text{size}_i$ becomes sufficiently small
- Suppose we contract as soon as $\alpha_i < 1/2$
- What is the problem?

Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:




Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...

- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:





Table contraction – the bad case

- Let n be an exact power of 2
- Suppose the first $n/2 + 1$ operations are of the type INSERT
- Suppose the last $n/2 - 1$ operations form the sequence:

DELETE, DELETE, INSERT, INSERT,
DELETE, DELETE, INSERT, INSERT, ...


- The first two of these causes a contraction, the next two an expansion, the next two a contraction, and so on:



- This is slow, even on average


Modifying the data structure

- We modify our approach by contracting when a DELETE operation causes the load factor α_i to become less than $1/4$




Modifying the data structure

- We modify our approach by contracting when a DELETE operation causes the load factor α_i to become less than $1/4$




Modifying the data structure

- We modify our approach by contracting when a DELETE operation causes the load factor α_i to become less than $1/4$



Modifying the data structure

- We modify our approach by contracting when a DELETE operation causes the load factor α_i to become less than $1/4$



- We expand in the same way as before, i.e., when an INSERT operation is applied to a full table T
- The modification avoids the problem on the previous slide
- Using the potential method, we show that the average time per



operation is $O(1)$


Wow!

The new potential function supporting INSERT and DELETE


- We pick the following potential function:

$$\Phi(D_i) = \begin{cases} 2\text{num}_i - \text{size}_i & \text{if } \alpha_i \geq 1/2 \\ \text{size}_i/2 - \text{num}_i & \text{if } \alpha_i < 1/2 \end{cases}$$

- Plotting the right-hand side as a function of num_i in the range $[\frac{1}{4}\text{size}_i, \text{size}_i]$ (in this range, size_i stays the same):



High-level proof that amortized cost is $\Theta(1)$



- Just before a table expansion/contraction, Φ has value num_i
- Just afterwards, the table is roughly half full and so the value of Φ drops to roughly 0
- This drop in potential pays for the expansion/contraction
- When no expansion/contraction occurs, the amortized cost is at most 3 since the slope of Φ has absolute value at most 2
- We now give a detailed proof of this

The potential function is valid

- Φ is a valid potential function, i.e., $\Phi(D_0) = 0$ and $\Phi(D_i) \geq 0$ for all i (see plot on previous slide)
- Hence,

$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i$$

- We will show that $\hat{c}_i \leq 3$ for $i = 1, \dots, n$ so that

$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i \leq 3n = O(n)$$


- This will show that the amortized cost of each update is $O(n)/n = O(1)$

Amortized cost of INSERT

- Consider the i th operation and assume that it is INSERT
- Assume first $\alpha_{i-1} \geq 1/2$ and $\alpha_i \geq 1/2$
- Then $\Phi(D_{i-1}) = 2\text{num}_{i-1} - \text{size}_{i-1}$ and $\Phi(D_i) = 2\text{num}_i - \text{size}_i$
- As shown earlier (when we only allowed insertions),
 $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) \leq 3$


Amortized cost of INSERT

- Now assume that not both $\alpha_{i-1} \geq 1/2$ and $\alpha_i \geq 1/2$
- Then $\alpha_{i-1} < 1/2$ and $\alpha_i \leq 1/2$ so table expansion cannot occur
- We have $\Phi(D_i) - \Phi(D_{i-1}) = -1$ so $\hat{c}_i = 1 - 1 = 0$:




Amortized cost of DELETE

- Next, assume that the i th operation is DELETE
- If no table contraction occurs, we have $c_i = 1$
- Thus, $\hat{c}_i \leq 2$ since $\Phi(D_i) - \Phi(D_{i-1}) \leq 1$:



Amortized cost of DELETE

- Next, assume that the i th operation is DELETE
- If no table contraction occurs, we have $c_i = 1$
- Thus, $\hat{c}_i \leq 2$ since $\Phi(D_i) - \Phi(D_{i-1}) \leq 1$:




Amortized cost of DELETE

- Now, assume that table contraction occurs
- In this case, one item is deleted and num_i items are moved to the contracted array
- Thus, the actual cost is $c_i = \text{num}_i + 1$
- If $\text{num}_i = 0$ then since $\text{size}_{i-1} \in \{1, 2, 4\}$, we have
 $\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) = 1 + 0 - \Phi(D_{i-1}) \leq 1$
- Now, assume $\text{num}_i > 0$
- Then $\alpha_{i-1} = 1/4 < 1/2$ and $\alpha_i < 1/2$ so we have
 $\Phi(D_i) = \text{size}_i/2 - \text{num}_i$ and $\Phi(D_{i-1}) = \text{size}_{i-1}/2 - \text{num}_{i-1}$
- The amortized cost is then

$$\begin{aligned}\hat{c}_i &= c_i + \Phi(D_i) - \Phi(D_{i-1}) \\ &= (\text{num}_i + 1) + (\text{size}_i/2 - \text{num}_i) - (\text{size}_{i-1}/2 - \text{num}_{i-1})\end{aligned}$$

Amortized cost of DELETE

- Illustration of table contraction:



- We express everything in terms of num_i :

$$\begin{aligned}\hat{c}_i &= c_i + \Phi(D_i) - \Phi(D_{i-1}) \\ &= num_i + 1 + (size_i/2 - num_i) - (size_{i-1}/2 - num_{i-1}) \\ &= num_i + 1 + ((num_i + 1) - num_i) \\ &\quad - (2(num_i + 1) - (num_i + 1)) \\ &= 1\end{aligned}$$

Bounding the average cost for dynamic table

- We have shown that $\hat{c}_i \leq 3$ for each operation i
- We can thus bound the total actual cost by

$$\sum_{i=1}^n c_i \leq \sum_{i=1}^n \hat{c}_i \leq 3n$$

- The average running time per operation is thus $O(n)/n = O(1)$

Plan for the lecture on February 27

- Fibonacci Heaps
- We will apply the potential method to analyze the performance of this data structure